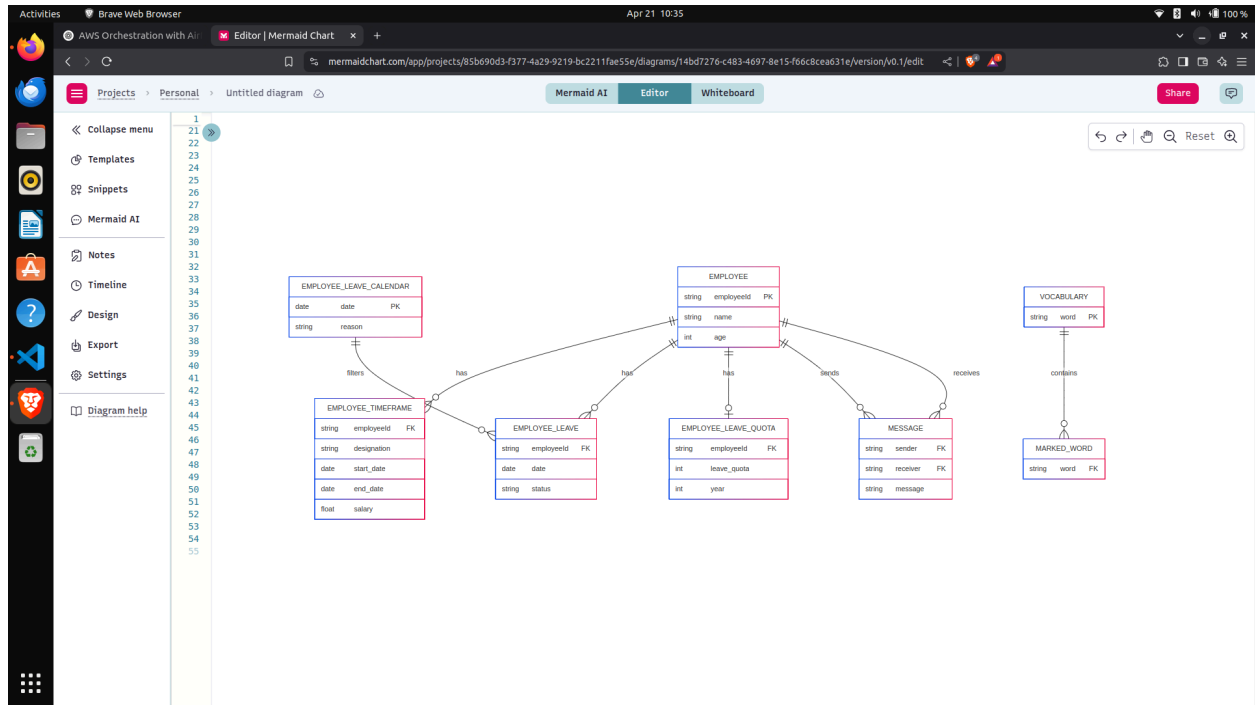


# PROJECT

## ER DIAGRAM



## erDiagram

```
EMPLOYEE {
    string employeeId PK
    string name
    int age
}
```

```
EMPLOYEE_TIMEFRAME {
    string employeeId FK
    string designation
    date start_date
    date end_date
    float salary
}
```

```
EMPLOYEE_LEAVE_QUOTA {
    string employeeId FK
    int leave_quota
    int year
}
```

# PROJECT

```
EMPLOYEE_LEAVE_CALENDAR {  
    date date PK  
    string reason  
}
```

```
EMPLOYEE_LEAVE {  
    string employeeId FK  
    date date  
    string status  
}
```

```
MESSAGE {  
    string sender FK  
    string receiver FK  
    string message  
}
```

```
VOCABULARY {  
    string word PK  
}
```

```
MARKED_WORD {  
    string word FK  
}
```

```
EMPLOYEE ||--o{ EMPLOYEE_TIMEFRAME : has  
EMPLOYEE ||--o{ EMPLOYEE_LEAVE : has  
EMPLOYEE ||--o| EMPLOYEE_LEAVE_QUOTA : has  
EMPLOYEE_LEAVE_CALENDAR ||--o{ EMPLOYEE_LEAVE : filters  
EMPLOYEE ||--o{ MESSAGE : sends  
EMPLOYEE ||--o{ MESSAGE : receives  
VOCABULARY ||--o{ MARKED_WORD : contains
```

# PROJECT

## STEP 1 CODE

```
import sys
import boto3
from datetime import datetime
from awsglue.context import GlueContext
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from pyspark.sql.functions import col

# Spark + Glue setup
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session

# Config
today = datetime.utcnow().strftime("%Y-%m-%d")
bucket = "poc-bootcamp-capstone-group1"
raw_prefix = "poc-bootcamp-group1-bronze/emp_data_qus1/raw/"
processed_output_path =
f"s3://{bucket}/poc-bootcamp-group1-bronze/emp_data_qus1/processed/data_pr
ocessed_{today}.csv"
silver_output_path =
f"s3://{bucket}/poc-bootcamp-group1-silver/employee_data/"

s3 = boto3.client('s3')

# Step 1: List raw CSV files
raw_files = s3.list_objects_v2(Bucket=bucket,
Prefix=raw_prefix).get("Contents", [])
csv_files = [obj["Key"] for obj in raw_files if
obj["Key"].endswith(".csv")]

if not csv_files:
    print("⚠️ No CSV files found in raw folder.")
    sys.exit(0)

print(f"📄 Found {len(csv_files)} CSV files.")

# Step 2: Merge and save raw files to processed zone
raw_paths = [f"s3://{bucket}/{key}" for key in csv_files]
```

# PROJECT

```
df_raw = spark.read.option("header", "true").csv(raw_paths)

# Write merged raw as CSV into processed folder (coalesced to 1 file)
df_raw.coalesce(1).write.mode("overwrite").option("header",
"true").csv(processed_output_path)
print(f"📦 Raw merged CSV copied to: {processed_output_path}")

# Delete raw files after backup
for key in csv_files:
    s3.delete_object(Bucket=bucket, Key=key)
    print(f"🗑 Deleted raw file: {key}")

# Step 3: Read processed file back for transformation
df_processed = spark.read.option("header",
"true").csv(processed_output_path)

# Step 4: Clean and transform
df_cleaned = df_processed.select(
    col("emp_id").cast("string"),
    col("age").cast("int"),
    col("name").cast("string")
).dropna().filter(col("age") > 0).dropDuplicates(["emp_id", "age",
"name"])

# Step 5: Write to silver zone (partitioned by folder, NOT ingestion
column)
df_cleaned.write.mode("append").parquet(silver_output_path)
print(f"✅ Processed & clean data written to: {silver_output_path}")

print("🎉 Glue job completed successfully.")
```

# PROJECT

## STEP 2

### PYSPARK CODE

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, from_unixtime, to_date, when,
row_number, lit, min as min_
from pyspark.sql.window import Window
import os
import shutil

# === Spark Session ===
spark = SparkSession.builder.appName("SCD2 Incremental").getOrCreate()
spark.sparkContext.setLogLevel("ERROR")

# === Folder paths ===
RAW_DIR = '/home/himanshu/Learning/bootcamp_project/QUS_02/data/raw'
STAGING_DIR =
'/home/himanshu/Learning/bootcamp_project/QUS_02/data/processed'
SILVER_PATH =
'/home/himanshu/Learning/bootcamp_project/QUS_02/data/silver/employees_scd
2.parquet'
PROCESSED_RAW_DIR =
'/home/himanshu/Learning/bootcamp_project/QUS_02/data/processed_raw'

os.makedirs(STAGING_DIR, exist_ok=True)

# === Step 1: Read raw file ===
raw_files = sorted([f for f in os.listdir(RAW_DIR) if f.endswith('.csv')])
if not raw_files:
    print("🚫 No raw file to process. Exiting.")
    exit()

file_name = raw_files[0]
raw_path = os.path.join(RAW_DIR, file_name)
print(f"\n📁 Processing file: {file_name}")

# === Step 2: Load and preprocess new file ===
df = spark.read.option("header", True).csv(raw_path)
df = df.withColumn("start_date",
to_date(from_unixtime(col("start_date").cast("long"))))
```

# PROJECT

```
df = df.withColumn("end_date",
to_date(from_unixtime(col("end_date").cast("long"))))
df = df.withColumn("salary", col("salary").cast("double"))
print("\n📄 Raw DataFrame after timestamp conversion:")
df.show()

# === Step 3: Deduplicate within file ===
w1 = Window.partitionBy("emp_id", "start_date",
"end_date").orderBy(col("salary").desc())
df = df.withColumn("rn1", row_number().over(w1)).filter(col("rn1") ==
1).drop("rn1")
print("\n✂️ After deduplication within raw data:")
df.show()

# === Step 4: Mark ACTIVE based on latest salary/null end ===
df = df.withColumn("status", when(col("end_date").isNull(),
"ACTIVE").otherwise("INACTIVE"))
print("\n✅ After status marking (ACTIVE/INACTIVE):")
df.show()

# === Step 5: Save processed version to staging
staging_path = os.path.join(STAGING_DIR, file_name.replace(".csv",
".parquet"))
df.write.mode("overwrite").parquet(staging_path)
print(f"\n📁 Saved processed file to staging at: {staging_path}")

# === Step 6: IF silver does not exist – First-time load
if not os.path.exists(SILVER_PATH):
    print("\n🆕 First-time load – writing data directly to silver.")
    df.write.mode("overwrite").parquet(SILVER_PATH)
    print("\n🟡 Silver Table (Initial Load):")
    df.show()
else:
    # === Step 7: SCD2 CONTINUATION LOGIC
    print("\n🔄 Silver table exists – continuing with SCD2 logic.")
    silver_df = spark.read.parquet(SILVER_PATH)
    processed_df = spark.read.parquet(staging_path)

    print("\n📁 Silver Table:")
    silver_df.show()
```

# PROJECT

```
print("\n📁 Processed DataFrame:")
processed_df.show()

active_df = silver_df.filter(col("status") == "ACTIVE")
inactive_df = silver_df.filter(col("status") == "INACTIVE")
print("\n🟢 ACTIVE Records:")
active_df.show()
print("\n🔴 INACTIVE Records:")
inactive_df.show()

# === Step 8: Get min start_date from new incoming rows
continuity_dates =
processed_df.groupBy("emp_id").agg(min_("start_date").alias("new_start_date"))
print("\n📅 Continuity Dates:")
continuity_dates.show()

# === Step 9: Update previous ACTIVE rows to INACTIVE
updated_old = active_df.alias("old") \
    .join(continuity_dates.alias("new"), "emp_id") \
    .filter(col("new.new_start_date") >= col("old.start_date")) \
    .withColumn("end_date", col("new.new_start_date")) \
    .withColumn("status", lit("INACTIVE")) \
    .select("old.emp_id", "old.designation", "old.start_date",
"end_date", "old.salary", "status")
print("\n✏️ Updated old ACTIVE rows:")
updated_old.show()

# === Step 10: Keep untouched active rows
untouched_active = active_df.join(continuity_dates.select("emp_id"),
"emp_id", "left_anti")
print("\n✅ Unchanged ACTIVE rows:")
untouched_active.show()

# === Step 11: Get truly new employees
new_emps = processed_df.join(silver_df.select("emp_id").distinct(),
"emp_id", "left_anti")
print("\n🆕 Truly new employees (not seen before):")
new_emps.show()
```

# PROJECT

```
# === Step 12: Merge final result
columns = ["emp_id", "designation", "start_date", "end_date", "salary",
"status"]

final_df = inactive_df.select(columns) \
    .union(updated_old.select(columns)) \
    .union(untouched_active.select(columns)) \
    .union(processed_df.select(columns)) \
    .union(new_emps.select(columns)) \
    .dropDuplicates(columns)

print("\n📊 Final merged Silver Table to write:")
final_df.show()

# === Step 13: Write final output to silver
final_df.write.mode("overwrite").parquet(SILVER_PATH)
print("\n💿 Final data written to Silver!")

# === Step 14: Cleanup: Move raw to processed_raw, delete staging
new_raw_path = os.path.join(PROCESSED_RAW_DIR, file_name)
shutil.move(raw_path, new_raw_path)
print(f"\n📦 Moved raw file to archive: {new_raw_path}")

shutil.rmtree(staging_path)
print(f"\n🧹 Cleaned staging folder: {staging_path}")

print(f"\n✅ Done. Silver updated and raw file archived: {file_name}")
```



# PROJECT

## GLUE JOB

```
from awsglue.utils import getResolvedOptions
from awsglue.context import GlueContext
from pyspark.context import SparkContext
from pyspark.sql import SparkSession, functions as F
from pyspark.sql.window import Window
from py4j.protocol import Py4JJavaError
from pyspark.sql.types import StructType, StructField, StringType,
DateType, DoubleType
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, from_unixtime, to_date, when
import pyspark.sql.functions as F
import sys
import boto3
import re

# === Glue Setup ===
args = getResolvedOptions(sys.argv, ['JOB_NAME'])
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session

# === Define Schema Manually ===
manual_schema = StructType([
    StructField("emp_id", StringType(), True),
    StructField("designation", StringType(), True),
    StructField("start_date", DateType(), True),
    StructField("end_date", DateType(), True),
    StructField("salary", DoubleType(), True),
    StructField("status", StringType(), True)
])

# === S3 Paths ===
RAW_PATH =
"s3://poc-bootcamp-capstone-group1/poc-bootcamp-group1-bronze/emp_timefram
e_data_qus2/raw/"
STAGING_PATH =
"s3://poc-bootcamp-capstone-group1/poc-bootcamp-group1-bronze/emp_timefram
e_data_qus2/processed/"
```

# PROJECT

```
PROCESSED_RAW_PATH =
"s3://poc-bootcamp-capstone-group1/poc-bootcamp-group1-bronze/emp_timefram
e_data_gus2/processed_raw/"
SILVER_PATH =
"s3://poc-bootcamp-capstone-group1/poc-bootcamp-group1-silver/emp-timefram
e-data/"

# === Step 1: Read raw CSV data
print(f"\n📂 Reading raw data from: {RAW_PATH}")
df = spark.read.option("header", True).csv(RAW_PATH)

# === Step 2: Cast salary and convert timestamps
df = df.withColumn("salary", F.col("salary").cast("double"))
df = df.withColumn("start_date",
F.to_date(F.from_unixtime(F.col("start_date").cast("long"))))
df = df.withColumn("end_date",
F.to_date(F.from_unixtime(F.col("end_date").cast("long"))))
print(f"\n📄 Raw DataFrame after cleaning and casting:")
df.show()

# === Step 3: Deduplicate
w1 = Window.partitionBy("emp_id", "start_date",
"end_date").orderBy(F.col("salary").desc())
df = df.withColumn("rn1", F.row_number().over(w1)).filter(F.col("rn1") ==
1).drop("rn1")

# === Step 4: Mark ACTIVE
df = df.withColumn("status", when(col("end_date").isNull(),
"ACTIVE").otherwise("INACTIVE"))
print(f"\n✅ After status marking (ACTIVE/INACTIVE):")

# === Step 5: Save to staging
df.write.mode("overwrite").parquet(STAGING_PATH)
print(f"\n💾 Saved processed file to staging at: {STAGING_PATH}")

# === Step 6: Check if silver path exists
s3 = boto3.client("s3")
bucket = "poc-bootcamp-capstone-group1"
key_prefix = "poc-bootcamp-group1-silver/emp-timeframe-data/"
response = s3.list_objects_v2(Bucket=bucket, Prefix=key_prefix)
```

# PROJECT

```
silver_exists = "Contents" in response

if not silver_exists:
    print("NEW First-time load – writing to silver.")
    df.write.mode("overwrite").parquet(SILVER_PATH)
else:
    print("\n🔄 Silver table exists – continuing with SCD2 logic.")

    try:
        silver_df = spark.read.schema(manual_schema).parquet(SILVER_PATH)
    except Py4JJavaError as e:
        if "UNABLE_TO_INFER_SCHEMA" in str(e):
            print("⚠️ Silver path is empty. Creating empty DataFrame.")
            silver_df = spark.createDataFrame([], schema=manual_schema)
        else:
            raise

    try:
        processed_df =
spark.read.schema(manual_schema).parquet(STAGING_PATH)
    except Py4JJavaError as e:
        if "UNABLE_TO_INFER_SCHEMA" in str(e):
            print("⚠️ Staging path is empty. Creating empty DataFrame.")
            processed_df = spark.createDataFrame([], schema=manual_schema)
        else:
            raise

    active_df = silver_df.filter(F.col("status") == "ACTIVE")
    inactive_df = silver_df.filter(F.col("status") == "INACTIVE")

    continuity_dates =
processed_df.groupBy("emp_id").agg(F.min("start_date").alias("new_start_da
te"))

    updated_old = active_df.alias("old") \
        .join(continuity_dates.alias("new"), "emp_id") \
        .filter(F.col("new.new_start_date") >= F.col("old.start_date")) \
        .withColumn("end_date", F.col("new.new_start_date")) \
        .withColumn("status", F.lit("INACTIVE")) \
```

# PROJECT

```
.select("old.emp_id", "old.designation", "old.start_date",
"end_date", "old.salary", "status")

untouched_active = active_df.join(continuity_dates.select("emp_id"),
"emp_id", "left_anti")

new_emps = processed_df.join(silver_df.select("emp_id").distinct(),
"emp_id", "left_anti")

columns = ["emp_id", "designation", "start_date", "end_date",
"salary", "status"]
final_df = inactive_df.select(columns) \
.union(updated_old.select(columns)) \
.union(untouched_active.select(columns)) \
.union(processed_df.select(columns)) \
.union(new_emps.select(columns)) \
.dropDuplicates(columns)

final_df.write.mode("overwrite").parquet(SILVER_PATH)
print("\n👁️ Final data written to Silver!")

# === Step 7: Move raw → processed_raw, delete only files from processed
def move_raw_files_to_archive(src_bucket, src_prefix, dst_prefix):
    print(f"\n📦 Moving raw files from s3://{src_bucket}/{src_prefix} to
s3://{src_bucket}/{dst_prefix}")
    response = s3.list_objects_v2(Bucket=src_bucket, Prefix=src_prefix)
    if "Contents" in response:
        for obj in response['Contents']:
            key = obj['Key']
            if not key.endswith('/'):
                copy_source = {'Bucket': src_bucket, 'Key': key}
                new_key = key.replace(src_prefix, dst_prefix, 1)
                s3.copy_object(Bucket=src_bucket, CopySource=copy_source,
Key=new_key)

                s3.delete_object(Bucket=src_bucket, Key=key)
                print(f"    - Moved: {key} → {new_key}")
    else:
        print("⚠️ No files in raw folder.")

def delete_files_only(bucket, prefix):
```

# PROJECT

```
print(f"\n🔪 Deleting files in s3://{bucket}/{prefix}")
response = s3.list_objects_v2(Bucket=bucket, Prefix=prefix)
if "Contents" in response:
    files = [obj['Key'] for obj in response['Contents'] if not
obj['Key'].endswith('/')]
    if files:
        s3.delete_objects(Bucket=bucket, Delete={'Objects': [{'Key':
k} for k in files]})
        for key in files:
            print(f"    - Deleted: {key}")
    else:
        print("⚠️ No file objects found.")
    else:
        print("⚠️ Nothing found under prefix.")

def extract_bucket_prefix(s3_path):
    match = re.match(r's3://([^/]+)/(.+)', s3_path)
    return match.groups() if match else (None, None)

raw_bucket, raw_prefix = extract_bucket_prefix(RAW_PATH)
staging_bucket, staging_prefix = extract_bucket_prefix(STAGING_PATH)
_, processed_raw_prefix = extract_bucket_prefix(PROCESSED_RAW_PATH)

if raw_bucket:
    move_raw_files_to_archive(raw_bucket, raw_prefix,
processed_raw_prefix)

if staging_bucket:
    delete_files_only(staging_bucket, staging_prefix)

print(f"\n✅ Glue Job Finished Successfully!")
```

# PROJECT

```
# Start Zookeeper
cd ~/kafka
bin/zookeeper-server-start.sh config/zookeeper.properties
```

```
# In a new terminal session, start Kafka
cd ~/kafka
bin/kafka-server-start.sh config/server.properties
```

```
# Create a topic named 'test-topic'
bin/kafka-topics.sh --create --topic test-topic --bootstrap-server localhost:9092 --partitions 1
--replication-factor 1
```

```
# List available topics
bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

```
# Start a producer to send messages
bin/kafka-console-producer.sh --topic test-topic --bootstrap-server localhost:9092
```

```
# In a new terminal session, start a consumer to receive messages
bin/kafka-console-consumer.sh --topic test-topic --from-beginning --bootstrap-server
localhost:9092
```

Producer running

```
(airflow_venv) ubuntu@ip-172-31-6-153:~$ python
/home/ubuntu/kafka_codes/KAFKA_QUS/producer.py
```

Consumer running

```
(airflow_venv) ubuntu@ip-172-31-6-153:~$ ~/spark/bin/spark-submit --master local[*] --jars
/home/ubuntu/postgresql-42.7.5.jar --packages
org.apache.spark:spark-sql-kafka-0-10_2.12:3.3.0,org.apache.kafka:kafka-clients:3.3.2
/home/ubuntu/kafka_codes/KAFKA_QUS/consumer_spark.py
```

Cooldown running

```
(airflow_venv) ubuntu@ip-172-31-6-153:~$ python
/home/ubuntu/kafka_codes/KAFKA_QUS/cooldown_strike_level_spark.py
```

# PROJECT

```
pkill -f producer.py  
pkill -f consumer_spark.py  
pkill -f cooldown_strike_level_spark.py
```

```
rm -f /tmp/80_leave_report_*.flag  
rm -f /tmp/leave_quota_done_*.flag
```

# PROJECT

[Spark DataFrame] → [Pandas DataFrame] → [CSV Buffer using StringIO] → [PostgreSQL COPY]

↑                    ↑                    ↑                    ↑

Massive data processing    Converts for COPY (Pandas supports it)    Acts like a file without disk I/O    Fast bulk insert

## Why Use Spark → Pandas → CSV → PostgreSQL?

### ◆ Reason: Spark can't `INSERT` into PostgreSQL efficiently

Spark can write to PostgreSQL via `.jdbc()` — but **only in full overwrite or append mode**, not row-by-row control. You needed:

1. **Precise DELETE** of old rows (via `psycopg2`)
2. **Efficient INSERT** of new rows — but PostgreSQL doesn't support high-performance inserts directly via `.jdbc()` for bulk rows.

That's why you:

- Used **Spark** for all data processing and transformation
- Then **converted Spark → Pandas → CSV buffer**
- And used **PostgreSQL's** `COPY FROM STDIN` for fast bulk insert

## Why `COPY FROM STDIN` is better?

Method	Speed	Control	Use Case
<code>.jdbc().write()</code>	Slower	None	Overwrite or append entire table
<code>INSERT INTO ... VALUES</code>	Very slow	Yes	Fine-grained control, but slow
<code>COPY FROM STDIN</code> (via <code>psycopg2</code> )	Very fast	Yes	Bulk insert millions of rows efficiently

`COPY FROM STDIN` works like loading from a file — but uses an in-memory `StringIO` buffer.






# PROJECT

Great question. Here's a clear and practical comparison to explain **why AWS Glue is more suitable than EMR** for your project.

## ✔ Why You Used AWS Glue Instead of EMR

Reason	AWS Glue Advantage	Why EMR Falls Short for Your Use Case 
1. Fully managed	No need to provision or manage clusters — serverless Spark	You need to configure and maintain EC2-based Spark/Hadoop clusters
2. Cost-efficient for ETL	Pay only per job execution + short duration = low cost	EMR runs on EC2; costs accumulate even if the cluster is idle
3. Optimized for ETL	Built-in connectors for S3, PostgreSQL, Redshift, JDBC, etc.	EMR supports these, but setup and integration take more work
4. Easy to maintain	Auto-handles scaling, retries, metadata (via Glue Data Catalog)	You handle all failure recovery, scaling, and setup
5. Native integration	Easily integrates with boto3, Athena, Lake Formation, and more	Integration is possible, but not as seamless
6. Trigger-based execution	Supports event, schedule, and dependency-based job triggers	You would need to set this up manually via Step Functions or Airflow
7. Ideal for small-to-medium pipelines	Perfect for jobs that run in minutes (like yours)	EMR is more suitable for long-running or massively parallel jobs



# PROJECT


## When EMR Would Be a Better Fit

Use EMR **instead of Glue** if:

- You need **custom big data engines** (like Hive, Presto, Flink, or HBase)
- You run **long-running Spark streaming** or batch jobs over TBs of data
- You require **fine-grained control** over hardware, memory, and cluster scaling
- You're doing **real-time processing** or **ML model training** at scale

## Your Project Profile

Feature	Characteristic
Data size	Small to moderate (CSV files, PostgreSQL writes)
Job frequency	Daily or event-based
Latency requirement	Low (completes in minutes)
User management effort	You want minimal overhead

 **Glue is a perfect match** for this scenario.



# PROJECT

## ✔ What Already Makes Your Pipeline Fault-Tolerant

Feature	Status	Why it helps	📄
Try-Except blocks	✔ Present	Catches runtime errors and logs stack trace for debugging	
Glue Job Commit	✔ Present	Ensures the Glue job marks itself complete only when successful	
Deduplication & null checks	✔ Present	Prevents data corruption from duplicates or nulls	
S3 and PostgreSQL separation	✔ Good design	Avoids tight coupling of storage and compute	
JDBC write modes ( <code>overwrite</code> )	⚠ Used safely	Works for single-target writes, but can overwrite all data if misused	
S3 report duplication check	✔ Present	Avoids redundant writes and reruns	

## ⚠ Where Fault Tolerance Is Missing / Needs Improvement

Missing Feature	Impact	📄
✗ Retries on failure	If a job fails due to a transient issue (e.g., network), it won't retry	
✗ Atomic PostgreSQL writes	<code>.jdbc()</code> writes are not transaction-controlled — can result in partial writes	
✗ Idempotency for data writes	Some <code>.write()</code> operations can duplicate or overwrite previous good data	
✗ Audit logs / checkpoints	There's no tracking of what data was ingested, when, or by which run	
✗ Upserts/Merge for final tables	Current overwrite/append logic may lead to duplicate entries	
⚠ Failure notifications	There's no alert system for failed jobs (like SNS, email, or CloudWatch Alarms)	

# PROJECT

## Recommendations to Make Your Pipeline Fully Fault-Tolerant

### 1. Enable Glue Job Retries

- Set `MaxRetries` in the job configuration or use workflows to auto-retry failed steps.

### 2. Use Atomic Writes or Staging Tables

- Instead of direct `.overwrite()` on final tables, write to a temp table and use SQL `MERGE` to update.
- Commit PostgreSQL transactions manually (via `psycopg2`) for large writes.

### 3. Idempotent Design

- Include a `run_id` or `ingest_batch_id` to avoid duplicate processing.
- Use `.dropDuplicates()` or `row_number()` in all staging logic.

### 4. Logging & Audit

- Write logs or checkpoints to a Glue table or S3 log file with `job_run_id`, `timestamp`, and status.


### 5. Monitoring & Alerts

- Add:
  - CloudWatch Alarms for failed Glue jobs
  - SNS alerts on exceptions
  - Logging to S3 for audit trails



# PROJECT

## ✓ 1. Where You Use Optimizations (Well Done)

Optimization	Applied?	Notes	
Broadcast join for small tables	✓	You used <code>broadcast()</code> when joining <code>leave_df</code> and <code>holiday_df</code> , which is ideal	
Filtering early (predicate pushdown)	✓	You filter <code>status</code> , <code>year</code> , and <code>date</code> before joins or writes (good practice)	
<code>.dropDuplicates()</code> and <code>.distinct()</code>	✓	Reduces write I/O and ensures data sanity	
Read with Schema	✓	You define schema for CSVs (prevents full scan and incorrect type inference)	
Use of <code>coalesce(1)</code> for small writes	✓	Ensures single output file during S3 saves — useful for small datasets	
Working day calendar logic	✓	Smart use of <code>dayofweek</code> and anti-joins to filter weekends/holidays efficiently	
Avoiding unnecessary columns	✓	You often <code>.select()</code> only what you need before aggregation or writing	

# PROJECT

## ❌ 2. Where Optimizations Are Missing or Skipped

Skipped Optimization	Impact
Partitioning when reading/writing large data	Reads entire tables instead of using filters like <code>pushDownPredicate</code> or <code>partitionColumn</code> with <code>.jdbc()</code>
No caching for reused DataFrames	Expensive joins or filters (like on leave data) are repeated across steps
No <code>.persist()</code> before multiple actions	E.g., multiple <code>.count()</code> and <code>.write()</code> on same <code>DataFrame</code> triggers recomputation
<code>.jdbc()</code> writes are row-by-row	JDBC writes can be slow if not batched or written via <code>COPY</code> using <code>psycopg2</code> for large data
No <code>bucketBy()</code> or <code>partitionBy()</code> on writes	Output tables aren't optimized for downstream queries
Data catalog partitioning (Glue Tables)	Not leveraged for faster S3 scans via Athena or Glue crawlers
No compression when writing to S3	You're writing raw CSVs — compressed formats like Parquet/ORC would save space and I/O

# PROJECT

## 3. How to Improve Optimization

- **For Reads (Especially JDBC)**

- Use:

```
python
.option("fetchsize", "10000")
.option("partitionColumn", "emp_id")
.option("lowerBound", "1")
.option("upperBound", "999999")
.option("numPartitions", "10")
```

This parallelizes JDBC reads.

- **For Writes**

- Avoid `.jdbc(mode="overwrite")` for big tables. Instead:
  - Use **staging table** → then `MERGE INTO`
  - Use `COPY FROM STDIN` via `psycpg2` for **bulk inserts**

- **DataFrame Caching**

- Use `.cache()` or `.persist()` on intermediate DataFrames accessed more than once (like filtered `leave_df`)

- **Partitioning / Bucketing**

- Use `.repartition("year")` or `.partitionBy("year")` when writing large S3 datasets

- **Format Optimization**

- Write to **Parquet** instead of CSV for S3:

```
python
df.write.option("compression", "snappy").parquet(path)
```

- **Reduce Actions on Large DataFrames**

- Avoid calling `.count()`, `.collect()` unless necessary

- **Use Glue Job Bookmarks (if doing incremental loads)**


- This allows Glue to **only pick new data** since last successful job

## 4. Optimization Strategy for Your Use Case

Pipeline Area	Optimization Goal	Recommendation
Leave Analysis (8%, 80%)	Low latency + efficient filters	Use caching, pushdown filters, avoid recompute
S3 writes	Reduce cost + increase scan speed	Use Parquet + <code>partitionBy("year")</code>
PostgreSQL I/O	Improve throughput	Use batching or <code>COPY</code> with <code>psycpg2</code>
Report generation	Avoid recomputation for multiple employees	Cache filtered and joined DataFrame before <code>.collect()</code>
SCD2 updates	Efficient incremental updates	Avoid full overwrite — use <code>MERGE INTO</code>

# PROJECT

## Final Scorecard: Optimization Maturity

Area	Score (✓=Good, ⚠=Some, ✗=Needs work)	
Read efficiency	✓	
Join strategies	✓	
Memory management	⚠	
Write performance	✗	
File formats	⚠	
Incremental logic	⚠	
Reuse / caching	✗	
Compression	✗	

### Bonus: Smart One-Liners (Use These in Answers)

- “We used `left_anti` joins to exclude holidays dynamically during working day filtering.”
- “To prevent duplicate reports on S3, we compare keys before generating new TXT files.”
- “For high-volume writes, I’d prefer `COPY FROM STDIN` using `psycopg2` over `.jdbc()`.”
- “We used `row_number()` and `window.partitionBy()` for deduplicating leave records.”
- “I tuned `.jdbc()` reads using `partitionColumn`, `lowerBound`, and `numPartitions`.”




# PROJECT

## What is Predicate Pushdown?


**Predicate pushdown** is an optimization technique where **filters** (`WHERE` clauses) are **pushed down to the data source level**, so **only necessary rows** are loaded into Spark from a data source like PostgreSQL, S3 (Parquet), or JDBC.

## Why It's Important

Without predicate pushdown:

- Spark loads **entire table/data file** into memory, then applies filters.
-  Slower, more memory-heavy

With predicate pushdown:

- Only filtered data is fetched from the **source**.
-  Faster reads, reduced I/O, more efficient memory use

# PROJECT

## ✓ How You Can Use It

### • In JDBC Reads

```
python Copy Edit

spark.read.jdbc(
    url=jdbc_url,
    table="(SELECT * FROM employee_db WHERE year = 2024) AS sub",
    properties=props
)
```

or better:

```
python Copy Edit

df = spark.read.jdbc(
    url=jdbc_url,
    table="employee_db",
    predicates=["year = 2024"],
    properties=db_properties
)
```

- Here, Spark sends `WHERE year = 2024` to PostgreSQL, reducing load

### • With Parquet / ORC on S3

If you:

- Use **columnar formats** like Parquet or ORC
- Apply `df.filter(col("year") == 2024)` **early**
- Glue or Spark can **read only the relevant row groups**

### • With CSV/JSON → ⚠ Limited



Predicate pushdown **does not work well** with plain CSV or JSON.

- Entire file still gets scanned.

# PROJECT

## How It Applies to Your Project

In your pipeline:

-  **Good:** You apply filters **early** (e.g., `status == 'ACTIVE'`, `date >= X`) before joins or writes
-  **Missed:** You could improve JDBC reads like this:

**Before (less optimal):**

python

 Copy

 Edit

```
df = spark.read.jdbc(url=jdbc_url, table="leave_data", properties=props)
df = df.filter(col("year") == CURRENT_YEAR)
```

**After (optimized):**

python

 Copy

 Edit

```
df = spark.read.jdbc(
    url=jdbc_url,
    table="leave_data",
    predicates=[f"year = {CURRENT_YEAR}"],
    properties=props
)
```

# PROJECT

## What is Adaptive Query Execution (AQE)?

Adaptive Query Execution (AQE) is a runtime optimization feature in Spark that **dynamically adjusts the execution plan** of a query **based on actual data statistics** gathered **during execution**.

✅ Introduced in **Spark 3.0+**

## Why Do We Need AQE?


Spark's Catalyst optimizer generates a **static execution plan before running**, based on **estimated statistics** (like table size, partition size, etc.).

But:

- Estimates can be wrong
- Joins can be inefficient
- Shuffles can be skewed

✅ **AQE fixes this** by adapting the plan **at runtime**, when actual data stats are known.

## Key Features of AQE

Feature	Description	
1. Dynamic Join Selection	Switches join strategy at runtime (e.g., from sort-merge to broadcast)	
2. Skew Join Handling	Detects skew and splits large partitions across reducers	
3. Dynamically Coalescing Shuffle Partitions	Reduces number of output partitions based on actual shuffle size	
4. Dynamically Switch Partition Pruning	Skips irrelevant partitions even during query runtime	

## ✅ How to Enable AQE

python


 Copy

 Edit

```
spark.conf.set("spark.sql.adaptive.enabled", "true")
```

# PROJECT

## `repartition()` vs `coalesce()` – Quick Summary

Feature	<code>repartition()</code>	<code>coalesce()</code> 
Purpose	Increases or decreases number of partitions	Primarily <b>reduces</b> number of partitions
Shuffle	✅ Full <b>shuffle</b>	⚠️ Avoids shuffle when reducing partitions
Use Case	When you want <b>even distribution</b> or increase partitions	When you want to <b>reduce partitions</b> efficiently
Performance Impact	Costly due to full data movement	Fast (no shuffle), but can create skew
Typical Use	Before joins, <code>groupBy</code> , wide shuffles	Before writing to disk (e.g., to S3/DB)

## What is Watermarking in Spark?

**Watermarking** is a technique used in **Structured Streaming** to **handle late-arriving data** and define **event time-based windows** with a **tolerance period** for lateness.

## Why Watermarking is Needed?

In streaming data:

- Data doesn't always arrive in order.
- Some records come **late** (due to network delays, retries, etc.)
- You need a way to **allow some lateness**, but also **garbage collect** old state.

✅ **Watermarking** balances this:

- Waits for **late events up to a threshold** (e.g., 10 minutes)
- Drops data that arrives **after that threshold**

# PROJECT